

210503H I.P.D.D.Rajapaksha

## 2 Learning from data


### 2.1. : Data generation.

```
import numpy as np
import matplotlib . pyplot as plt
from sklearn . model_selection import train_test_split
from sklearn . linear_model import LinearRegression
# Generate 100 samples
n_samples = 100
# Generate X values ( uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)
# Generate epsilon values ( normally distributed with mean 0 and
standard deviation 15)
epsilon = np . random . normal (0 , 15 , n_samples )
# Generate Y values using the model  $Y = 3 + 3X + \epsilon$ 
Y = 3 + 3 * X + epsilon [: , np . newaxis ]
```

### 2.2. : Data visualization.

```
r=np.random.randint(104)
# Split the data into training and test sets (80% train,20% test)
X_train, X_test, Y_train, Y_test = train_test_split(X,
Y,test_size=0.2,random_state=r)
# Plot the data points

plt.figure(figsize=(10, 6))
plt.scatter(X_train, Y_train, alpha=1,
marker='o',color='red',label='Training_Data')
plt.scatter(X_test, Y_test, alpha=1,
marker='s',color='blue',label='Testing_Data')
plt.show()
```



The reason the training and testing datasets differ in each run is due to the use of the `np.random.randint(104)` function, which generates a random integer each time it's called. This integer acts as a random seed for the `train_test_split` function, determining how the data is shuffled and then divided into training and testing sets.

When you set a specific value for the random state in the train-test data splitting, it guarantees that the same data points will be included in the training and testing sets every time the code is run, ensuring consistency. However, in this case, we are using `r=np.random.randint(104)` to generate a different random state for each run. Because `r` changes randomly every time the code is executed, it results in different data points being selected for training and testing, which is why we observe different scatter plots each time we run the code.

## Extra part to compare difference

```
import numpy as np
import matplotlib.pyplot as plt
#from sklearn.model_selection import train_test_split

# Generate four different random states
random_states = [np.random.randint(104) for _ in range(4)]

# Create a 2x2 grid of subplots
fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# Loop over the random states and axes to create the plots
for i, ax in enumerate(axes.flat):
    # Split the data using the current random state
```

```

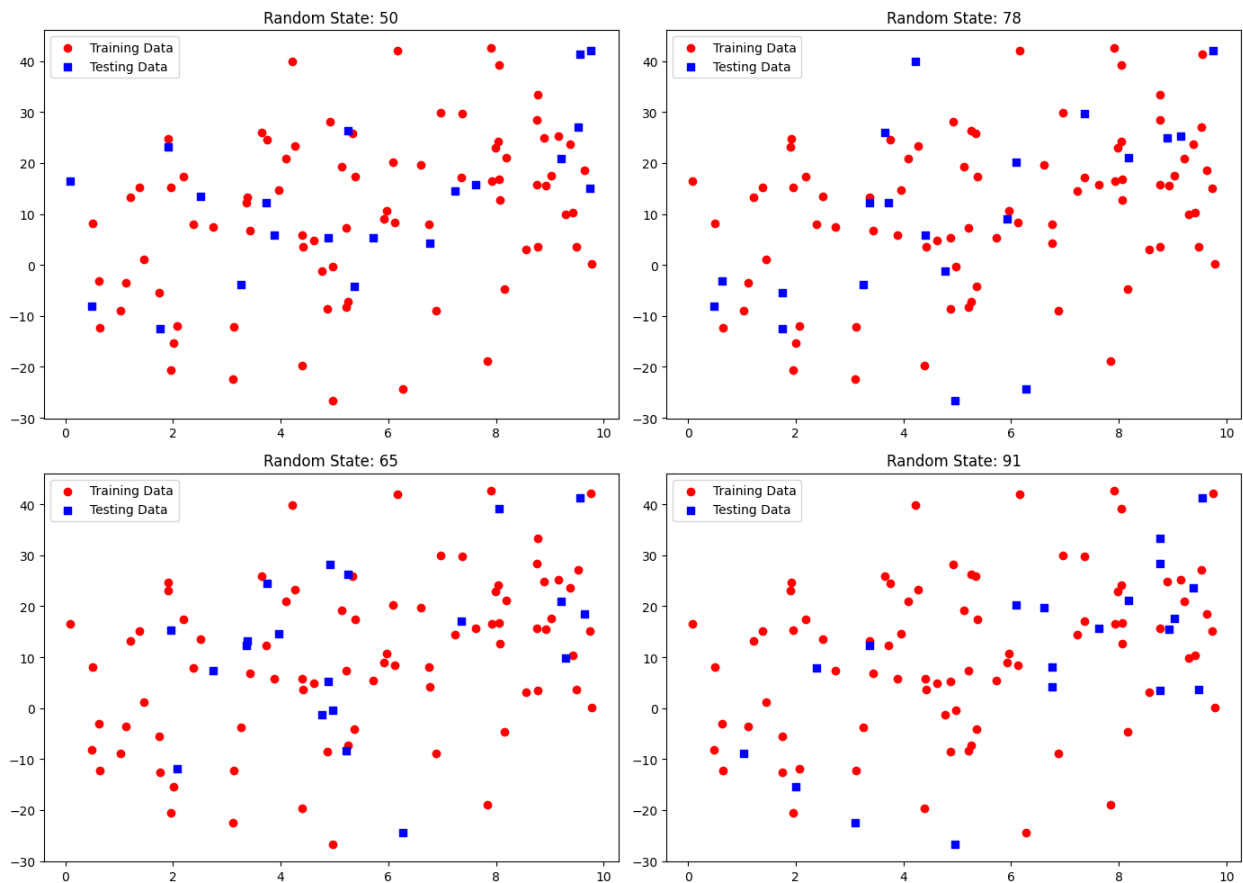
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state=random_states[i])

# Plot the data points
ax.scatter(X_train, Y_train, alpha=1, marker='o', color='red',
label='Training Data')
ax.scatter(X_test, Y_test, alpha=1, marker='s', color='blue',
label='Testing Data')

# Add title and legend to each subplot
ax.set_title(f'Random State: {random_states[i]}')
ax.legend()

# Display the grid of images
plt.tight_layout()
plt.show()


```



When state change testing and training data set will change. But overall same data set was chosen for both.

## 2.3. Linear regression


```
for i in range(10):# Plotting 10 different instances
    X_train, X_test, Y_train, Y_test = train_test_split(X,Y,
test_size=0.2, random_state=np.random.randint(104))
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```



The linear regression model differs from one instance to another because the training dataset changes each time due to the random splitting of data using a randomly generated random state. Each time the `train_test_split` function is called with a new random state, different training data is selected. Since the model is trained on different data sets in each iteration, the model parameters (such as slope and intercept) change, resulting in different linear regression models for each iteration.

## 2.4.Repeat 3 with n\_sample = 10,000

```
# Generate 10,000 samples
n_samples = 10000
# Generate X values ( uniformly distributed between 0 and 10)
X = 10 * np.random.rand(n_samples, 1)
# Generate epsilon values ( normally distributed with mean 0 and
standard deviation 15)
epsilon = np . random . normal (0 , 15 , n_samples )
# Generate Y values using the model  $Y = 3 + 3X +$ 
epsilon
Y = 3 + 2 * X + epsilon [: , np . newaxis ]
for i in range(10):# Plotting 10 different instances
    X_train, X_test, Y_train, Y_test = train_test_split(X,Y,
test_size=0.2, random_state=np.random.randint(104))
    model = LinearRegression()
    model.fit(X_train, Y_train)
    Y_pred_train = model.predict(X_train)
    plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')
plt.xlabel('X')
plt.ylabel('Y')
plt.legend()
plt.show()
```



When the number of data samples is increased to 10,000, the variation in the linear regression models across different iterations will generally decrease. This happens because larger datasets provide more information, leading to more stable and consistent model training. With 10,000 samples, the training data is more representative of the overall data distribution, reducing the impact of random variations caused by different random states in the data splitting process. Consequently, the model parameters (such as slope and intercept) become more stable, resulting in similar linear regression lines across different iterations. In contrast, with only 100 samples, the random splits lead to greater variability in the training data, causing more fluctuations in the model parameters and the resulting regression lines. (When we increase the sample size population mean will equal to sample mean)

## 3 Linear regression on real world data

### 3.1. Loading data

Installing missing package

```
!pip install ucimlrepo

Collecting ucimlrepo
  Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: pandas>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2.1.4)
Requirement already satisfied: certifi>=2020.12.5 in
/usr/local/lib/python3.10/dist-packages (from ucimlrepo) (2024.8.30)
Requirement already satisfied: numpy<2,>=1.22.4 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2024.1)
Requirement already satisfied: tzdata>=2022.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->ucimlrepo) (2024.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.16.0)
Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7

# If package not installed, install it using pip install ucimlrepo
from ucimlrepo import fetch_ucirepo
# fetch dataset
infrared_thermography_temperature = fetch_ucirepo(id=925)
```

```

# data (as pandas dataframes)
X = infrared_thermography_temperature.data.features
y = infrared_thermography_temperature.data.targets
# metadata
print(infrared_thermography_temperature.metadata)
# variable information
print(infrared_thermography_temperature.variables)

{'uci_id': 925, 'name': 'Infrared Thermography Temperature',
 'repository_url':
 'https://archive.ics.uci.edu/dataset/925/infrared+thermography+tempera
 ture+dataset', 'data_url':
 'https://archive.ics.uci.edu/static/public/925/data.csv', 'abstract':
 'The Infrared Thermography Temperature Dataset contains temperatures
 read from various locations of inferred images about patients, with
 the addition of oral temperatures measured for each individual. The 33
 features consist of gender, age, ethnicity, ambient temperature,
 humidity, distance, and other temperature readings from the thermal
 images. The dataset is intended to be used in a regression task to
 predict the oral temperature using the environment information as well
 as the thermal image readings.', 'area': 'Health and Medicine',
 'tasks': ['Regression'], 'characteristics': ['Tabular'],
 'num_instances': 1020, 'num_features': 33, 'feature_types': ['Real',
 'Categorical'], 'demographics': ['Gender', 'Age', 'Ethnicity'],
 'target_col': ['aveOralF', 'aveOralM'], 'index_col': ['SubjectID'],
 'has_missing_values': 'no', 'missing_values_symbol': None,
 'year_of_dataset_creation': 2021, 'last_updated': 'Tue Dec 12 2023',
 'dataset_doi': '10.13026/9ay4-2c37', 'creators': ['Quanzeng Wang',
 'Yangling Zhou', 'Pejman Ghassemi', 'David McBride', 'J. Casamento',
 'T. Pfefer', 'Quanzeng Wang', 'Yangling Zhou', 'Pejman Ghassemi',
 'David McBride', 'J. Casamento', 'T. Pfefer'], 'intro_paper':
 {'title': 'Infrared Thermography for Measuring Elevated Body
 Temperature: Clinical Accuracy, Calibration, and Evaluation',
 'authors': 'Quanzeng Wang, Yangling Zhou, Pejman Ghassemi, David
 McBride, J. Casamento, T. Pfefer', 'published_in': 'Italian National
 Conference on Sensors', 'year': 2021, 'url':
 'https://www.semanticscholar.org/paper/443b9932d295ca3a014e7d874b4bd77
 a33a276bd', 'doi': None}, 'additional_info': {'summary': None,
 'purpose': None, 'funded_by': None, 'instances_represent': None,
 'recommended_data_splits': None, 'sensitive_data': None,
 'preprocessing_description': None, 'variable_info': '- gender\n- age\n-
 ethnicity\n- ambient temperature\n- humidity\n- distance\n-
 temperature readings from the thermal images', 'citation': None},
 'external_url':
 'https://physionet.org/content/face-oral-temp-data/1.0.0/'}

   name      role      type demographic \
0  SubjectID    ID  Categorical      None
1    aveOralF  Target  Continuous      None
2    aveOralM  Target  Continuous      None
3      Gender  Feature  Categorical  Gender

```

4	Age	Feature	Categorical	Age
5	Ethnicity	Feature	Categorical	Ethnicity
6	T_atm	Feature	Continuous	None
7	Humidity	Feature	Continuous	None
8	Distance	Feature	Continuous	None
9	T_offset1	Feature	Continuous	None
10	Max1R13_1	Feature	Continuous	None
11	Max1L13_1	Feature	Continuous	None
12	aveAllR13_1	Feature	Continuous	None
13	aveAllL13_1	Feature	Continuous	None
14	T_RC1	Feature	Continuous	None
15	T_RC_Dry1	Feature	Continuous	None
16	T_RC_Wet1	Feature	Continuous	None
17	T_RC_Max1	Feature	Continuous	None
18	T_LC1	Feature	Continuous	None
19	T_LC_Dry1	Feature	Continuous	None
20	T_LC_Wet1	Feature	Continuous	None
21	T_LC_Max1	Feature	Continuous	None
22	RCC1	Feature	Continuous	None
23	LCC1	Feature	Continuous	None
24	canthiMax1	Feature	Continuous	None
25	canthi4Max1	Feature	Continuous	None
26	T_FHCC1	Feature	Continuous	None
27	T_FHRC1	Feature	Continuous	None
28	T_FHLC1	Feature	Continuous	None
29	T_FHBC1	Feature	Continuous	None
30	T_FHTC1	Feature	Continuous	None
31	T_FH_Max1	Feature	Continuous	None
32	T_FHC_Max1	Feature	Continuous	None
33	T_Max1	Feature	Continuous	None
34	T_OR1	Feature	Continuous	None
35	T_OR_Max1	Feature	Continuous	None

description units		
missing_values		
0	Subject ID	None
no		
1	Oral temperature measured in fast mode	None
no		
2	Oral temperature measured in monitor mode	None
no		
3	Male or Female	None
no		
4	Age ranges in categories\n	None
no		
5	American Indian or Alaska Native, Asian, Black...	None
no		
6	Ambiant temperature	None
no		



7	Relative humidity	None
no		
8	Distance between the subjects and the IRTs.	None
no		
9	Temperature difference between the set and mea...	None
no		
10	Max value of a circle with diameter of 13 pixe...	None
no		
11	Max value of a circle with diameter of 13 pixe...	None
no		
12	Average value of a circle with diameter of 13 ...	None
no		
13	Average value of a circle with diameter of 13 ...	None
no		
14	Average temperature of the highest four pixels...	None
no		
15	Average temperature of the highest four pixels...	None
no		
16	Average temperature of the highest four pixels...	None
no		
17	Max value of a square of 24x24 pixels around t...	None
no		
18	Average temperature of the highest four pixels...	None
no		
19	Average temperature of the highest four pixels...	None
no		
20	Average temperature of the highest four pixels...	None
no		
21	Max value of a circle with diameter of 13 pixe...	None
no		
22	Average value of a square of 3x3 pixels center...	None
no		
23	Average value of a square of 3x3 pixels center...	None
no		
24	Max value in the extended canthi area	None
no		
25	Average temperature of the highest four pixels...	None
no		
26	Average value in the center point of forehead,...	None
no		
27	Average value in the right point of the forehe...	None
no		
28	Average value in the left point of the forehea...	None
no		
29	Average value in the bottom point of the foreh...	None
no		
30	Average value in the top point of the forehead...	None
no		
31	Maximum temperature within the extended forehe...	None

```

no
32 Max value in the center point of forehead, a s... None
no
33 Maximum temperature within the whole face region. None
no
34 Average temperature of the highest four pixels... None
no
35 Maximum temperature within the mouth region. None
no

```

## 3.2.

```

print(X.shape[1])#to cheak number of independent variables
print(y.shape[1])#to cheak number of dependent variables

33
2

```

numer of independet variables - 33

numer of dependet variables - 2

## 3.3.

```

print(infrared_thermography_temperature.metadata)
print(infrared_thermography_temperature.variables)

{'uci_id': 925, 'name': 'Infrared Thermography Temperature',
 'repository_url':
 'https://archive.ics.uci.edu/dataset/925/infrared+thermography+tempera
 ture+dataset', 'data_url':
 'https://archive.ics.uci.edu/static/public/925/data.csv', 'abstract':
 'The Infrared Thermography Temperature Dataset contains temperatures
 read from various locations of inferred images about patients, with
 the addition of oral temperatures measured for each individual. The 33
 features consist of gender, age, ethnicity, ambient temperature,
 humidity, distance, and other temperature readings from the thermal
 images. The dataset is intended to be used in a regression task to
 predict the oral temperature using the environment information as well
 as the thermal image readings.', 'area': 'Health and Medicine',
 'tasks': ['Regression'], 'characteristics': ['Tabular'],
 'num_instances': 1020, 'num_features': 33, 'feature_types': ['Real',
 'Categorical'], 'demographics': ['Gender', 'Age', 'Ethnicity'],
 'target_col': ['ave0ralF', 'ave0ralM'], 'index_col': ['SubjectID'],
 'has_missing_values': 'no', 'missing_values_symbol': None,
 'year_of_dataset_creation': 2021, 'last_updated': 'Tue Dec 12 2023',
 'dataset_doi': '10.13026/9ay4-2c37', 'creators': ['Quanzeng Wang',
 'Yangling Zhou', 'Pejman Ghassemi', 'David McBride', 'J. Casamento',
 'T. Pfefer', 'Quanzeng Wang', 'Yangling Zhou', 'Pejman Ghassemi',
 'David McBride', 'J. Casamento', 'T. Pfefer'], 'intro_paper':

```

```
{'title': 'Infrared Thermography for Measuring Elevated Body
Temperature: Clinical Accuracy, Calibration, and Evaluation',
'authors': 'Quanzeng Wang, Yangling Zhou, Pejman Ghassemi, David
McBride, J. Casamento, T. Pfefer', 'published_in': 'Italian National
Conference on Sensors', 'year': 2021, 'url':
'https://www.semanticscholar.org/paper/443b9932d295ca3a014e7d874b4bd77
a33a276bd', 'doi': None}, 'additional_info': {'summary': None,
'purpose': None, 'funded_by': None, 'instances_represent': None,
'recommended_data_splits': None, 'sensitive_data': None,
'preprocessing_description': None, 'variable_info': '- gender\n- age\n-
ethnicity\n- ambient temperature\n- humidity\n- distance\n-
temperature readings from the thermal images', 'citation': None},
'external_url':
'https://physionet.org/content/face-oral-temp-data/1.0.0/'}
```

	name	role	type	demographic \
0	SubjectID	ID	Categorical	None
1	ave0ralF	Target	Continuous	None
2	ave0ralM	Target	Continuous	None
3	Gender	Feature	Categorical	Gender
4	Age	Feature	Categorical	Age
5	Ethnicity	Feature	Categorical	Ethnicity
6	T_atm	Feature	Continuous	None
7	Humidity	Feature	Continuous	None
8	Distance	Feature	Continuous	None
9	T_offset1	Feature	Continuous	None
10	Max1R13_1	Feature	Continuous	None
11	Max1L13_1	Feature	Continuous	None
12	aveAllR13_1	Feature	Continuous	None
13	aveAllL13_1	Feature	Continuous	None
14	T_RC1	Feature	Continuous	None
15	T_RC_Dry1	Feature	Continuous	None
16	T_RC_Wet1	Feature	Continuous	None
17	T_RC_Max1	Feature	Continuous	None
18	T_LC1	Feature	Continuous	None
19	T_LC_Dry1	Feature	Continuous	None
20	T_LC_Wet1	Feature	Continuous	None
21	T_LC_Max1	Feature	Continuous	None
22	RCC1	Feature	Continuous	None
23	LCC1	Feature	Continuous	None
24	canthiMax1	Feature	Continuous	None
25	canthi4Max1	Feature	Continuous	None
26	T_FHCC1	Feature	Continuous	None
27	T_FHRC1	Feature	Continuous	None
28	T_FHLC1	Feature	Continuous	None
29	T_FHBC1	Feature	Continuous	None
30	T_FHTC1	Feature	Continuous	None
31	T_FH_Max1	Feature	Continuous	None
32	T_FHC_Max1	Feature	Continuous	None
33	T_Max1	Feature	Continuous	None

34	T_OR1	Feature	Continuous	None	
35	T_OR_Max1	Feature	Continuous	None	
					description units
missing_values					
0	Subject ID				None
no					
1	Oral temperature measured in fast mode				None
no					
2	Oral temperature measured in monitor mode				None
no					
3	Male or Female				None
no					
4	Age ranges in categories\n				None
no					
5	American Indian or Alaska Native, Asian, Black...				None
no					
6	Ambiant temperature				None
no					
7	Relative humidity				None
no					
8	Distance between the subjects and the IRTs.				None
no					
9	Temperature difference between the set and mea...				None
no					
10	Max value of a circle with diameter of 13 pixe...				None
no					
11	Max value of a circle with diameter of 13 pixe...				None
no					
12	Average value of a circle with diameter of 13 ...				None
no					
13	Average value of a circle with diameter of 13 ...				None
no					
14	Average temperature of the highest four pixels...				None
no					
15	Average temperature of the highest four pixels...				None
no					
16	Average temperature of the highest four pixels...				None
no					
17	Max value of a square of 24x24 pixels around t...				None
no					
18	Average temperature of the highest four pixels...				None
no					
19	Average temperature of the highest four pixels...				None
no					
20	Average temperature of the highest four pixels...				None
no					
21	Max value of a circle with diameter of 13 pixe...				None
no					

22	Average value of a square of 3x3 pixels center...	None
no		
23	Average value of a square of 3x3 pixels center...	None
no		
24	Max value in the extended canthi area	None
no		
25	Average temperature of the highest four pixels...	None
no		
26	Average value in the center point of forehead,...	None
no		
27	Average value in the right point of the forehe...	None
no		
28	Average value in the left point of the forehea...	None
no		
29	Average value in the bottom point of the foreh...	None
no		
30	Average value in the top point of the forehead...	None
no		
31	Maximum temperature within the extended forehe...	None
no		
32	Max value in the center point of forehead, a s...	None
no		
33	Maximum temperature within the whole face region.	None
no		
34	Average temperature of the highest four pixels...	None
no		
35	Maximum temperature within the mouth region.	None
no		

It is not possible to apply linear regression directly on this dataset because it contains categorical data, such as 'Gender', 'Age', and 'Ethnicity'. Linear regression requires numerical input, so categorical data must be converted into a numerical form to proceed.

#### Steps:

1. **One-Hot Encoding:** Convert the categorical features ('Gender', 'Age', 'Ethnicity') into numerical values using one-hot encoding. This will create binary columns for each category, making them suitable for linear regression.
2. **Feature Scaling:** After encoding, features are scaled appropriately.
3. **Handling Missing Values:** Check for any missing values in the dataset remove them or replce them with average values.
4. **Check for Multicollinearity:** Cheak highly correlated features. If present, combine or remove correlated features to avoid redundancy and improve model performance.

### 3.4.

The provided code snippet attempts to handle missing values by dropping them separately from X and y: # Drop rows with missing values from both X and y X = X.dropna() y = y.dropna()

The provided method is wrong because it removes missing values from X and y independently. This separate handling can cause the features and target labels to become misaligned. In a dataset, each row in the feature set X should correspond directly to a row in the target variable y. If rows are removed from X without removing the corresponding rows from y, or vice versa, the pairing of data points with their respective labels is disrupted.

To correctly handle the missing values, we first need to inspect the dataset to understand how many missing values are present. If only a few values are missing, it might be acceptable to drop the corresponding rows.

```
import pandas as pd
data_set = pd.concat([X, y], axis=1)

print(data_set.isnull().sum())
```

Gender	0
Age	0
Ethnicity	0
T_atm	0
Humidity	0
Distance	2
T_offset1	0
Max1R13_1	0
Max1L13_1	0
aveAllR13_1	0
aveAllL13_1	0
T_RC1	0
T_RC_Dry1	0
T_RC_Wet1	0
T_RC_Max1	0
T_LC1	0
T_LC_Dry1	0
T_LC_Wet1	0
T_LC_Max1	0
RCC1	0
LCC1	0
canthiMax1	0
canthi4Max1	0
T_FHCC1	0
T_FHRC1	0
T_FHLC1	0
T_FHBC1	0
T_FHTC1	0
T_FH_Max1	0
T_FHC_Max1	0

```
T_Max1      0
T_OR1       0
T_OR_Max1   0
ave0ra1F    0
ave0ra1M    0
dtype: int64
```

Output showing that only 'Distance' column has two missing values. As there are only two missing values in 'Distance', it's okay to remove these rows (with respect to total data set).

```
data_set = data_set.dropna(subset=['Distance'])
```

```
#data_set =data_set.dropna()
```

```
print(data_set.isnull().sum())
```

```
Gender      0
Age         0
Ethnicity   0
T_atm       0
Humidity    0
Distance    0
T_offset1   0
Max1R13_1   0
Max1L13_1   0
aveAllR13_1 0
aveAllL13_1 0
T_RC1       0
T_RC_Dry1   0
T_RC_Wet1   0
T_RC_Max1   0
T_LC1       0
T_LC_Dry1   0
T_LC_Wet1   0
T_LC_Max1   0
RCC1        0
LCC1        0
canthiMax1  0
canthi4Max1 0
T_FHCC1     0
T_FHRC1     0
T_FHLC1     0
T_FHBC1     0
T_FHTC1     0
T_FH_Max1   0
T_FHC_Max1  0
T_Max1      0
T_OR1       0
T_OR_Max1   0
```

```
ave0ra1F      0
ave0ra1M      0
dtype: int64
```

```
#import pandas as pd
#data_set = pd.concat([X, y], axis=1)
print(data_set)
print("-----")
data_set.head()
```

	Gender	Age	Ethnicity	T_atm	Humidity
Distance \					
0	Male	41-50	White	24.0	28.0
0.8					
1	Female	31-40	Black or African-American	24.0	26.0
0.8					
2	Female	21-30	White	24.0	26.0
0.8					
3	Female	21-30	Black or African-American	24.0	27.0
0.8					
4	Male	18-20	White	24.0	27.0
0.8					
...	...	...	...	...	...
...					
1015	Female	21-25	Asian	25.7	50.8
0.6					
1016	Female	21-25	White	25.7	50.8
0.6					
1017	Female	18-20	Black or African-American	28.0	24.3
0.6					
1018	Male	26-30	Hispanic/Latino	25.0	39.8
0.6					
1019	Female	18-20	White	23.8	45.6
0.6					

	T_offset1	Max1R13_1	Max1L13_1	aveAllR13_1	...	T_FHLC1
T_FHBC1 \						
0	0.7025	35.0300	35.3775	34.4000	...	33.3725
33.4925						
1	0.7800	34.5500	34.5200	33.9300	...	33.6775
33.9700						
2	0.8625	35.6525	35.5175	34.2775	...	34.6475
34.8200						
3	0.9300	35.2225	35.6125	34.3850	...	34.6550
34.3025						
4	0.8950	35.5450	35.6650	34.9100	...	34.3975
34.6700						
...	...	...	...	...	...	...
...						



1015	1.2225	35.6425	35.6525	34.8575	...	35.4000
35.1375						
1016	1.4675	35.9825	35.7575	35.4275	...	35.2200
35.2075						
1017	0.1300	36.4075	36.3400	35.8700	...	35.2275
35.3675						
1018	1.2450	35.8150	35.5250	34.2950	...	34.9250
34.7150						
1019	0.8675	35.7075	35.5825	34.8875	...	34.6700
34.2150						

	T_FHTC1	T_FH_Max1	T_FHC_Max1	T_Max1	T_OR1	T_OR_Max1
ave0ralF \						
0	33.0025	34.5300	34.0075	35.6925	35.6350	35.6525
36.85						
1	34.0025	34.6825	34.6600	35.1750	35.0925	35.1075
37.00						
2	34.6700	35.3450	35.2225	35.9125	35.8600	35.8850
37.20						
3	34.9175	35.6025	35.3150	35.7200	34.9650	34.9825
36.85						
4	33.8275	35.4175	35.3725	35.8950	35.5875	35.6175
36.80						
...	...	...	...	...	...	...
...						
1015	35.2750	35.8525	35.7475	36.0675	35.6775	35.7100
36.95						
1016	35.0700	35.7650	35.5525	36.5000	36.4525	36.4900
37.25						
1017	35.3425	36.3750	35.7100	36.5350	35.9650	35.9975
37.35						
1018	34.5950	35.4150	35.3100	35.8600	35.4150	35.4350
37.15						
1019	34.7100	35.1525	35.1175	35.9725	35.8900	35.9175
37.05						

	ave0ralM
0	36.59
1	37.19
2	37.34
3	37.09
4	37.04
...	...
1015	36.99
1016	37.19
1017	37.59
1018	37.29
1019	37.19

[1018 rows x 35 columns]

```
-----  
-  
  
{"type": "dataframe", "variable_name": "data_set"}
```

### 3.5.

select "aveOralM" as the dependent feature and 'T\_atm', 'Humidity', 'Distance', 'T\_FH\_Max1', 'age' as independent variables

```
# One-hot encode the 'Age' categorical variable  
age_encoded = pd.get_dummies(data_set['Age'], prefix='Age')  
  
# Add the one-hot encoded columns back to the dataset and drop the  
# original 'Age' column  
data_set = pd.concat([data_set, age_encoded], axis=1)  
data_set = data_set.drop('Age', axis=1)  
  
# Select the dependent and independent variables  
dependent_feature = data_set['aveOralM']  
independent_features = data_set[['T_atm', 'Humidity', 'Distance',  
                                'T_FH_Max1']] + list(age_encoded.columns)  
  
# Combine into a new DataFrame for modeling  
model_data = pd.concat([independent_features, dependent_feature],  
                        axis=1)  
  
# Display the prepared dataset  
print("Prepared dataset:")  
independent_features.head()
```

Prepared dataset:

```
{"summary": "{\n  \"name\": \"independent_features\",\n  \"rows\": 1018,\n  \"fields\": [\n    {\n      \"column\": \"T_atm\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 1.337616782238765,\n        \"min\": 20.2,\n        \"max\": 29.1,\n        \"num_unique_values\": 78,\n        \"samples\": [\n          26.1,\n          24.0,\n          25.4\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Humidity\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 13.070899796157311,\n        \"min\": 9.9,\n        \"max\": 61.2,\n        \"num_unique_values\": 353,\n        \"samples\": [\n          53.2,\n          10.6,\n          50.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Distance\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 2.4564860370861084,\n        \"min\": 0.54,\n        \"max\": 79.0,\n        \"num_unique_values\": 33,\n        \"samples\": [\n          0.57,\n          0.63,\n          0.67\n        ],\n        \"semantic_type\": \"\"\n      }\n    ]\n  }\n}
```



```
3    37.09
4    37.04
Name: ave0ralM, dtype: float64
```

## 3.6.

Split the data into training and testing sets with 80% of data points for training and 20% of data points for testing.

```
X_train,X_test, y_train, y_test =
train_test_split(independent_features, dependent_feature,
test_size=0.2, random_state=42)
```

## 3.7.

Train a linear regression model and estimate the coefficient corresponds to independent variables.

```
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

# Make predictions using the testing set
y_pred = model.predict(X_test)

# Creating the DataFrame with features and their coefficients
coeff_df = pd.DataFrame({'Feature': X_train.columns, 'Coefficient':
model.coef_})

# Install the tabulate module
!pip install tabulate

# Import the tabulate function
from tabulate import tabulate

# Displaying the DataFrame as a table
print("Intercept:", model.intercept_)
print("Estimated features vs coefficients:")
print(tabulate(coeff_df, headers='keys', tablefmt='pretty'))


Requirement already satisfied: tabulate in
/usr/local/lib/python3.10/dist-packages (0.9.0)
Intercept: 13.238798996804732
Estimated features vs coefficients:
+-----+-----+-----+-----+-----+
```

	Feature	Coefficient
0	T_atm	-0.055535893239226766
1	Humidity	0.003422392005267033
2	Distance	0.0026393537872089093
3	T_FH_Max1	0.7092038079633628
4	Age_18-20	-0.09521023788460654
5	Age_21-25	-0.10911840778186253
6	Age_21-30	-0.03372839990148858
7	Age_26-30	-0.08394109824533079
8	Age_31-40	-0.13698172540270473
9	Age_41-50	0.21089662330066358
10	Age_51-60	-0.08176010339345889
11	Age_>60	0.32984334930878834

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         'k--', lw=2, color='red')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs. Predicted Values')
plt.show()
```

<ipython-input-22-53a471df34c0>:3: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "k--" (-> color='k'). The keyword argument will take precedence.

```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         'k--', lw=2, color='red')
```



### 3.8.

By looking at the coefficient we can see T\_FH\_Max1 got the maximum value which mean T\_FH\_Max1 contributes highly for the dependent feature

```
# Plotting the coefficients
!pip install seaborn # install seaborn
import seaborn as sns
plt.figure(figsize=(10, 6))
sns.barplot(x='Coefficient', y='Feature', data=coeff_df,
palette='viridis')
plt.title('Feature Coefficients in Linear Regression Model')
plt.xlabel('Coefficient Value')
plt.ylabel('Feature')
plt.show()
```

```
Requirement already satisfied: seaborn in
/usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (1.26.4)
Requirement already satisfied: pandas>=1.2 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (2.1.4)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in
/usr/local/lib/python3.10/dist-packages (from seaborn) (3.7.1)
```

Requirement already satisfied: contourpy>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.0)

Requirement already satisfied: cycler>=0.10 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.53.1)

Requirement already satisfied: kiwisolver>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.5)

Requirement already satisfied: packaging>=20.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (24.1)

Requirement already satisfied: pillow>=6.2.0 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (9.4.0)

Requirement already satisfied: pyparsing>=2.3.1 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.1.4)

Requirement already satisfied: python-dateutil>=2.7 in  
/usr/local/lib/python3.10/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in  
/usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.1)


Requirement already satisfied: tzdata>=2022.1 in  
/usr/local/lib/python3.10/dist-packages (from pandas>=1.2->seaborn) (2024.1)

Requirement already satisfied: six>=1.5 in  
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)

<ipython-input-23-843087b86571>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Coefficient', y='Feature', data=coeff_df,  
palette='viridis')
```



### 3.9.

```
# Select the dependent and independent variables
dependent_feature_n = data_set['aveOralM']
independent_features_n = data_set[['T_OR1', 'T_OR_Max1', 'T_FHC_Max1',
'T_FH_Max1']] ]
# Combine into a new DataFrame for modeling
model_data_n = pd.concat([independent_features_n,
dependent_feature_n], axis=1)

# Display the prepared dataset
print("Prepared dataset:")
#independent_features_n.head()

X_train,X_test, y_train, y_test =
train_test_split(independent_features_n, dependent_feature_n,
test_size=0.2, random_state=42)

from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions using the testing set
y_pred = model.predict(X_test)
```



```
# Creating the DataFrame with features and their coefficients
coeff_df = pd.DataFrame({'Feature': X_train.columns, 'Coefficient':
model.coef_})
```

```
# Displaying the DataFrame as a table
print("Intercept:", model.intercept_)
print("Estimated features vs coefficients:")
print(tabulate(coeff_df, headers='keys', tablefmt='pretty'))
```

Prepared dataset:

Intercept: 6.79355629984887


Estimated features vs coefficients:

	Feature	Coefficient
0	T_OR1	0.20545776323994563
1	T_OR_Max1	0.34819684316002775
2	T_FHC_Max1	-0.08371846705362093
3	T_FH_Max1	0.376564342065323

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
'k--', lw=2, color='red')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs. Predicted Values')
plt.show()
```

<ipython-input-25-53a471df34c0>:3: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "k--" (-> color='k'). The keyword argument will take precedence.


```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
'k--', lw=2, color='red')
```




```
import matplotlib.pyplot as plt

# Assuming 'independent_features_n' is your DataFrame with independent
# features
# and 'dependent_feature_n' is your Series with the target variable
for col in independent_features_n.columns:
    plt.figure(figsize=(8, 6))
    plt.scatter(independent_features_n[col], dependent_feature_n,
alpha=0.6)
    plt.xlabel(col) # Label with the independent variable name
    plt.ylabel('Target Variable')
    plt.title(f'Scatter Plot of {col} vs. Target Variable')
    plt.show()
```


Scatter Plot of T\_OR1 vs. Target Variable




Scatter Plot of T\_OR\_Max1 vs. Target Variable



Scatter Plot of T\_FHC\_Max1 vs. Target Variable





```
yhat=model.predict(X_train)
d = independent_features_n.shape[1]
print('d:',d)

# Residual Sum of Squares (RSS)
RSS = np.sum((yhat - y_train)**2)
print('RSS=', RSS)

N=len(y_train)
#print('Number of Datapoints=',N)

# Residual Standard Error (RSE)
RSE = np.sqrt(1/(N-d-1)*RSS)
print('RSE=', RSE)

# Mean square error(MSE)
#predictions = lm.predict(X_train)

newX =
pd.DataFrame({"Constant":np.ones(len(X_train))}).join(pd.DataFrame(X_t
```

```

rain))
MSE = (sum((y_train-yhat)**2))/(len(newX)-len(newX.columns))
print('MSE=', MSE)

# Total Sum of Squares (TSS)
TSS = np.sum((y_train- np.mean(y_train))**2)
print('TSS=', TSS)

# R2

R2 = (TSS - RSS)/TSS
print('R² (from direct calculations)=', R2)

# Calculation of R2 using sklearn
R2 = model.score(X_train,y_train)

print('R² (from sklearn module)=', R2)

d: 4
RSS= 77.97449082857881
RSE= 0.31045739777400483
MSE= 0.09638379583260669
TSS= 223.63911855036855
R² (from direct calculations)= 0.6513378726673116
R² (from sklearn module)= 0.6513378726673116

from scipy.stats import t
import numpy as np
import scipy.stats as stats
#samples size
SN = len(X_train)

features = ['T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']
w_0 = model.intercept_
w_1 = model.coef_

NF = len(features)

#Calculate Standard error
standard_error = []
for feature in features:
    SE2 = RSS/(SN-NF-1) / np.sum((X_train[feature] -
np.mean(X_train[feature]))**2)
    standard_error.append(np.sqrt(SE2))

#calculate t values
t_values = []
for i in range(NF):
    t_values.append(w_1[i]/standard_error[i])

#calculate p-values

```

```

p_values = []
for i in range(NF):
    p_values.append(2 * (1 - stats.t.cdf(abs(t_values[i]), SN-NF-1)))

# Create a DataFrame to display the values in a table
results_df = pd.DataFrame({
    'Feature': features,
    'Standard Error': standard_error,
    't-value': t_values,
    'p-value': p_values
})

# Print the table
print(results_df)

# Save the table to a CSV file (optional)
results_df.to_csv('model_statistics.csv', index=False)

```

	Feature	Standard Error	t-value	p-value
0	T_OR1	0.019226	10.686208	0.00000
1	T_OR_Max1	0.019219	18.117535	0.00000
2	T_FHC_Max1	0.018810	-4.450810	0.00001
3	T_FH_Max1	0.020771	18.129689	0.00000