



Department of Electronic and Telecommunication  
Engineering

University of Moratuwa

**EN3150-Pattern Recognition**

**Assignment 01:**  
**Learning from data and related challenges**  
**and linear models for regression**

Rajapaksha I.P.D.D. - 210503H

# Contents

<b>1 Datapre-processing</b>	<b>3</b>
<b>2 Learning from data</b>	<b>4</b>
2.1 Question 2 : Reason to training and testing data is different in each run . . . . .	4
2.2 Question 3 : Linear regression model is different from one instance to other . . . . .	6
2.3 Question 4 : Increase the number of data samples to 10,000 . . . . .	6
<b>3 Linear regression on real world data</b>	<b>8</b>
3.1 Question 1 : Load the dataset . . . . .	8
3.2 Question 2 : Numer of independent variables and dependent variables in the data set . . . .	8
3.3 Question 3 : possibility to apply linear regression and steps before applying linear regression	8
3.4 Question 4 : NaN/missing values and correct approach . . . . .	9
3.5 Question 5 : Select 'aveOralm' as the dependent feature and 'T_atm', 'Humidity', 'Distance', 'T_FH_Max1', 'age' as independent variables . . . . .	10
3.6 Question 6 : Split the data into training and testing sets . . . . .	10
3.7 Question 7 : Train a linear regression model and estimate the coefficient corresponds to independent variables. . . . .	11
3.8 Question 8 : Highly contributes independent variable. . . . .	12
3.9 Question 9 : Select 'T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1' fea- tures as independent features . . . . .	13
3.10 Question 9 : Calculations . . . . .	14
3.11 Question 11 . . . . .	16
<b>4 Performance evaluation of Linear regression</b>	<b>16</b>
4.1 Question 2 : Residual standard error (RSE) for models A and B . . . . .	16
4.2 Question 3 : R-squared (R2) for models A and B . . . . .	16
4.3 Question 4 : Between RSE and R-squared which performance metric is more fair . . . . .	16
<b>5 Impact of Outliers on Linear Regression</b>	<b>17</b>
5.1 Question2 : when $a \rightarrow 0$ . . . . .	17
5.2 Question 3 : To minimize the influence of data points . . . . .	17

## 1. Datapre-processing

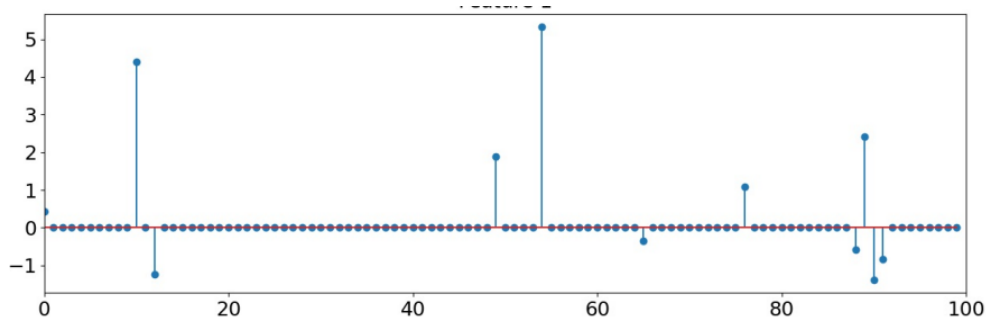


Figure 1: Feature 1 values of a dataset.

The values are mostly clustered around zero, with some occasional outliers reaching values between -1 and 5. Standard Scaling is suitable here because it will transform the feature to have a mean of 0 and a standard deviation of 1. This method handles outliers relatively well, preventing them from dominating the scaled values.

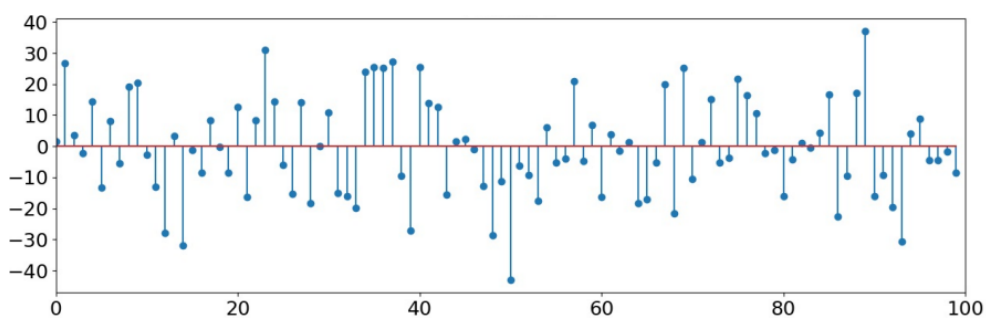


Figure 2: Feature 1 values of a dataset.

The feature values vary widely, ranging from approximately -40 to +40. The distribution is more spread out compared to Feature 1. Min-Max Scaling is suitable for Feature 2 because it scales the values to a [0, 1] range. This method is effective for maintaining the relationships between data points while minimizing the impact of outliers. It keeps the values within a defined range, preserving their relative distances.

## 2. Learning from data

### 2.1 Question 2 : Reason to training and testing data is different in each run

```
1 r=np.random.randint(104)
2 # Split the data into training and test sets (80% train,20% test)
3 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state
    =r)
4 # Plot the data points
5
6 plt.figure(figsize=(10, 6))
7 plt.scatter(X_train, Y_train, alpha=1, marker='o', color='red', label='Training Data'
    )
8 plt.scatter(X_test, Y_test, alpha=1, marker='s', color='blue', label='Testing Data')
9 plt.show()
```

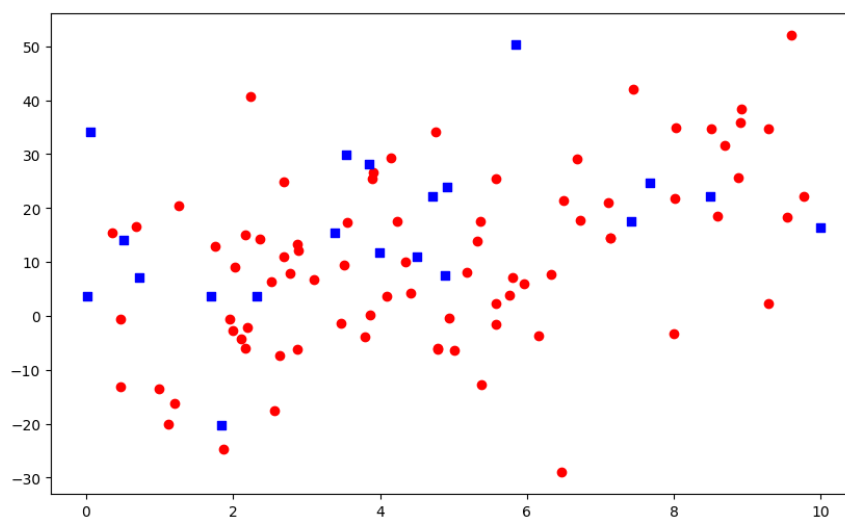


Figure 3: First Run

The reason the training and testing datasets differ in each run is due to the use of the `np.random.randint(104)` function, which generates a random integer each time it's called. This integer acts as a random seed for the `train_test_split` function, determining how the data is shuffled and then divided into training and testing sets.

When you set a specific value for the random state in the train-test data splitting, it guarantees that the same data points will be included in the training and testing sets every time the code is run, ensuring consistency. However, in this case, we are using `r=np.random.randint(104)` to generate a different random state for each run. Because `r` changes randomly every time the code is executed, it results in different data points being selected for training and testing, which is why we observe different scatter plots each time we run the code.

### Extra part to compare difference

```
10 import numpy as np
11 import matplotlib.pyplot as plt
12 #from sklearn.model_selection import train_test_split
13
14 # Generate four different random states
15 random_states = [np.random.randint(104) for _ in range(4)]
16
17 # Create a 2x2 grid of subplots
18 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
19
20 # Loop over the random states and axes to create the plots
21 for i, ax in enumerate(axes.flat):
22     # Split the data using the current random state
23     X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
24                                                         random_state=random_states[i])
25
26     # Plot the data points
27     ax.scatter(X_train, Y_train, alpha=1, marker='o', color='red', label='Training
28               Data')
29     ax.scatter(X_test, Y_test, alpha=1, marker='s', color='blue', label='Testing Data
30               ')
31
32     # Add title and legend to each subplot
33     ax.set_title(f'Random State: {random_states[i]}')
34     ax.legend()
35
36 # Display the grid of images
37 plt.tight_layout()
38 plt.show()
```

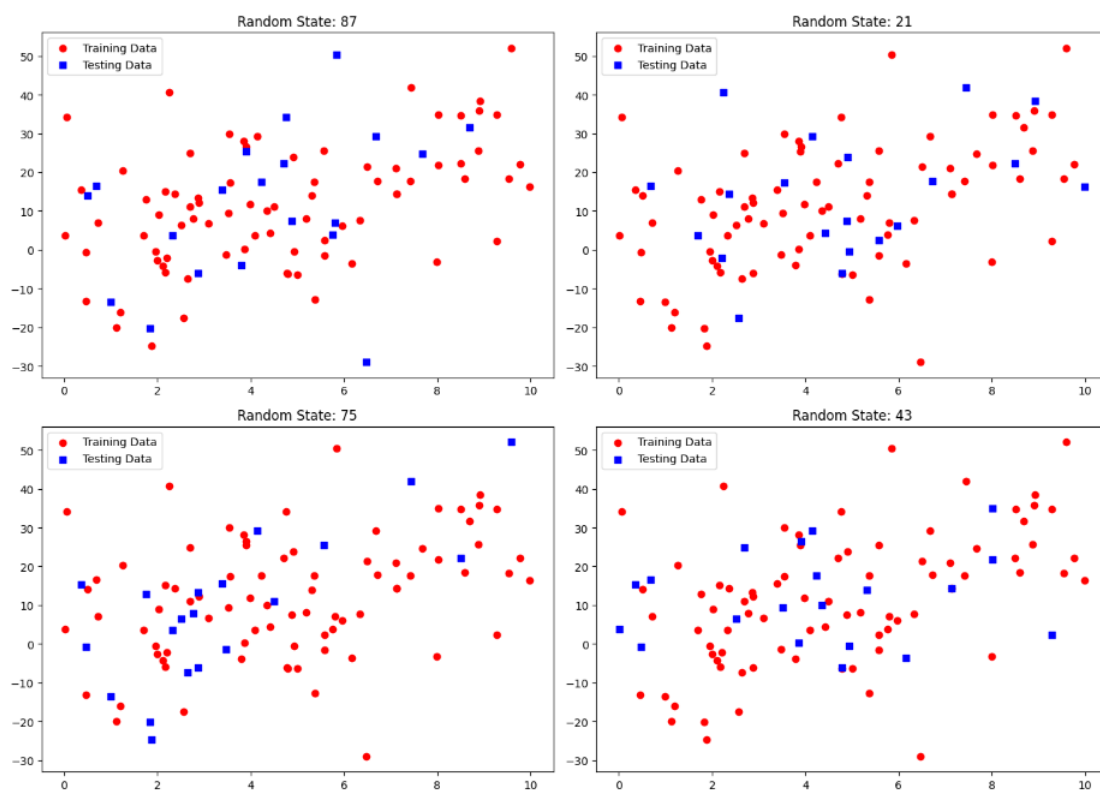


Figure 4: Code Output

When state change testing and training data set will change. But overall same data set was chosen for both.

## 2.2 Question 3 : Linear regression model is different from one instance to other

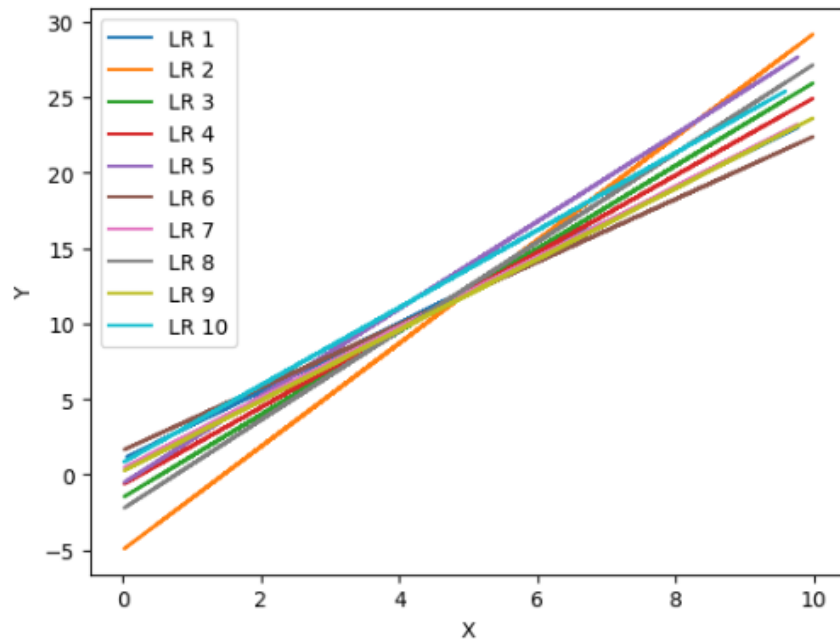


Figure 5: linear regression model

The linear regression model differs from one instance to another because the training dataset changes each time due to the random splitting of data using a randomly generated random state. Each time the `train_test_split` function is called with a new random state, different training data is selected. Since the model is trained on different data sets in each iteration, the model parameters (such as slope and intercept) change, resulting in different linear regression models for each iteration.

## 2.3 Question 4 : Increase the number of data samples to 10,000

```
37 # Generate 10,000 samples
38 n_samples = 10000
39 # Generate X values ( uniformly distributed between 0 and 10)
40 X = 10 * np.random.rand(n_samples, 1)
41 # Generate epsilon values ( normally distributed with mean 0 and standard deviation
42   15)
43 epsilon = np . random . normal (0 , 15 , n_samples )
44 # Generate Y values using the model Y = 3 + 3X +
45   epsilon
46 Y = 3 + 2 * X + epsilon [: , np . newaxis ]
47 for i in range(10):# Plotting 10 different instances
48     X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2,
49         random_state=np.random.randint(104))
50     model = LinearRegression()
51     model.fit(X_train, Y_train)
52     Y_pred_train = model.predict(X_train)
53     plt.plot(X_train, Y_pred_train, label=f'LR {i+1}')
54 plt.xlabel('X')
55 plt.ylabel('Y')
56 plt.legend()
57 plt.show()
```

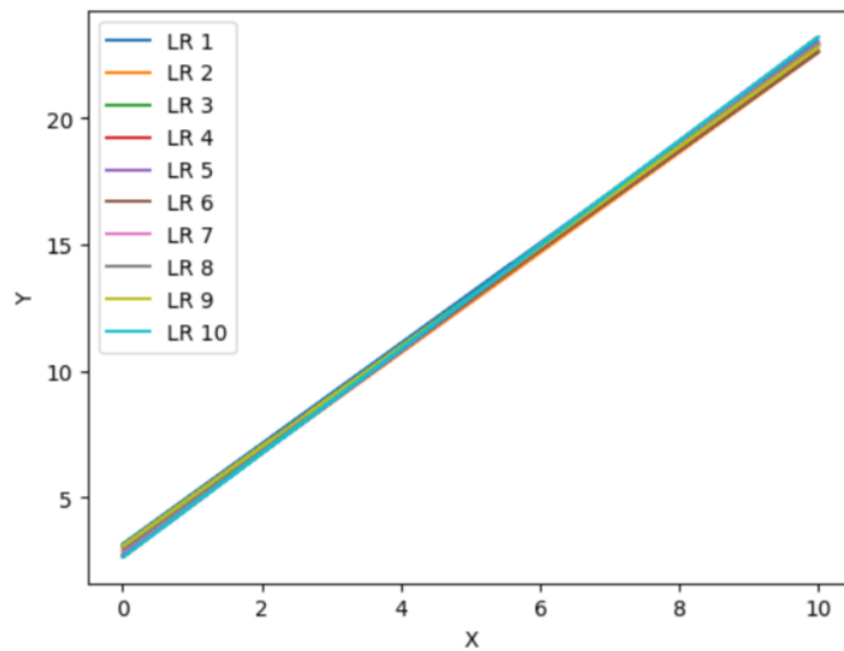


Figure 6: linear regression model with `n_sample=10000`

When the number of data samples is increased to 10,000, the variation in the linear regression models across different iterations will generally decrease. This happens because larger datasets provide more information, leading to more stable and consistent model training. With 10,000 samples, the training data is more representative of the overall data distribution, reducing the impact of random variations caused by different random states in the data splitting process. Consequently, the model parameters (such as slope and intercept) become more stable, resulting in similar linear regression lines across different iterations. In contrast, with only 100 samples, the random splits lead to greater variability in the training data, causing more fluctuations in the model parameters and the resulting regression lines. (When we increase the sample size population mean will equal to sample mean)

### 3. Linear regression on real world data

#### 3.1 Question 1 : Load the dataset

Installing missing package

```
1 !pip install ucimlrepo
```

#### 3.2 Question 2 : Numer of independent variables and dependent variables in the data set

```
1 print(X.shape[1]) #to cheak number of independent variables
2 print(y.shape[1]) #to cheak number of dependent variables
```

Number of independent variables = 33

Number of independent variables = 2

#### 3.3 Question 3 : possibility to apply linear regression and steps before applying linear regression

```
1 print(infrared_thermography_temperature.metadata)
2 print(infrared_thermography_temperature.variables)
```

	name	role	type	demographic
0	SubjectID	ID	Categorical	None
1	aveOralF	Target	Continuous	None
2	aveOralM	Target	Continuous	None
3	Gender	Feature	Categorical	Gender
4	Age	Feature	Categorical	Age
5	Ethnicity	Feature	Categorical	Ethnicity
6	T_atm	Feature	Continuous	None
7	Humidity	Feature	Continuous	None
8	Distance	Feature	Continuous	None
9	T_offset1	Feature	Continuous	None

Figure 7: Output

It is not possible to apply linear regression directly on this dataset because it contains categorical data, such as 'Gender', 'Age', and 'Ethnicity'. Linear regression requires numerical input, so categorical data must be converted into a numerical form to proceed.

##### Steps:

**1.One-Hot Encoding:** Convert the categorical features ('Gender', 'Age', 'Ethnicity') into numerical values using one-hot encoding. This will create binary columns for each category, making them suitable for linear regression.

**2.Feature Scaling:** After encoding, features are scaled appropriately.

**3.Handling Missing Values:** Check for any missing values in the dataset remove them or replce them with average values.

**4.Check for Multicollinearity:**Cheak highly correlated features. If present, combine or remove correlated features to avoid redundancy and improve model performance.



### 3.4 Question 4 : NaN/missing values and correct approach

The provided code snippet attempts to handle missing values by dropping them separately from X and y:

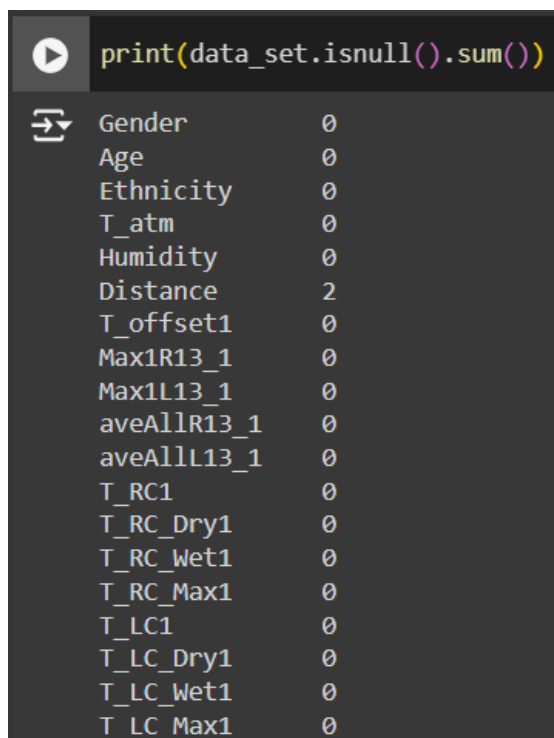
```
1 # Drop rows with missing values from both X and y
2 X = X.dropna()
3 y = y.dropna()
```

The provided method is wrong because it removes missing values from X and y independently. This separate handling can cause the features and target labels to become misaligned. In a dataset, each row in the feature set X should correspond directly to a row in the target variable y. If rows are removed from X without removing the corresponding rows from y, or vice versa, the pairing of data points with their respective labels is disrupted.

To correctly handle the missing values, we first need to inspect the dataset to understand how many missing values are present. If only a few values are missing, it might be acceptable to drop the corresponding rows.

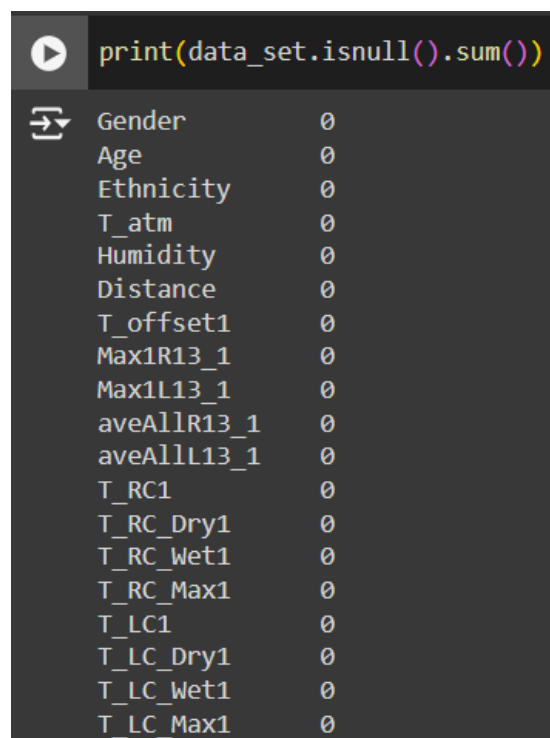
**Corrected code:**

```
1 import pandas as pd
2 data_set = pd.concat([X, y], axis=1)
3 data_set = data_set.dropna(subset=['Distance'])
```



Gender	0
Age	0
Ethnicity	0
T_atm	0
Humidity	0
Distance	2
T_offset1	0
Max1R13_1	0
Max1L13_1	0
aveAllR13_1	0
aveAllL13_1	0
T_RC1	0
T_RC_Dry1	0
T_RC_Wet1	0
T_RC_Max1	0
T_LC1	0
T_LC_Dry1	0
T_LC_Wet1	0
T_LC_Max1	0

(a) before run the code



Gender	0
Age	0
Ethnicity	0
T_atm	0
Humidity	0
Distance	0
T_offset1	0
Max1R13_1	0
Max1L13_1	0
aveAllR13_1	0
aveAllL13_1	0
T_RC1	0
T_RC_Dry1	0
T_RC_Wet1	0
T_RC_Max1	0
T_LC1	0
T_LC_Dry1	0
T_LC_Wet1	0
T_LC_Max1	0

(b) Caption for second image

Figure 8: after run the code

### 3.5 Question 5 : Select 'aveOralM' as the dependent feature and 'T\_atm', 'Humidity', 'Distance', 'T\_FH\_Max1', 'age' as independent variables

```
1 # One-hot encode the 'Age' categorical variable
2 age_encoded = pd.get_dummies(data_set['Age'], prefix='Age')
3
4 # Add the one-hot encoded columns back to the dataset and drop the original 'Age'
  column
5 data_set = pd.concat([data_set, age_encoded], axis=1)
6 data_set = data_set.drop('Age', axis=1)
7
8 # Select the dependent and independent variables
9 dependent_feature = data_set['aveOralM']
10 independent_features = data_set[['T_atm', 'Humidity', 'Distance', 'T_FH_Max1'] + list
   (age_encoded.columns)]
11
12 # Combine into a new DataFrame for modeling
13 model_data = pd.concat([independent_features, dependent_feature], axis=1)
14
15 # Display the prepared dataset
16 print("Prepared dataset:")
17 independent_features.head()
```

Prepared dataset:

	T_atm	Humidity	Distance	T_FH_Max1	Age_18-20	Age_21-25	Age_21-30	Age_26-30	Age_31-40	Age_41-50	Age_51-60	Age_>60
0	24.0	28.0	0.8	34.5300	False	False	False	False	False	True	False	False
1	24.0	26.0	0.8	34.6825	False	False	False	False	True	False	False	False
2	24.0	26.0	0.8	35.3450	False	False	True	False	False	False	False	False
3	24.0	27.0	0.8	35.6025	False	False	True	False	False	False	False	False
4	24.0	27.0	0.8	35.4175	True	False	False	False	False	False	False	False

Figure 9: Output

```
1 dependent_feature.head()
```

	aveOralM
0	36.59
1	37.19
2	37.34
3	37.09
4	37.04
dtype: float64	

Figure 10: Output

### 3.6 Question 6 : Split the data into training and testing sets

```
1 X_train,X_test, y_train, y_test = train_test_split(independent_features,
   dependent_feature, test_size=0.2, random_state=42)
```

### 3.7 Question 7 : Train a linear regression model and estimate the coefficient corresponds to independent variables.

```
1 from sklearn import datasets, linear_model
2 from sklearn.metrics import mean_squared_error, r2_score
3
4 model = LinearRegression()
5 model.fit(X_train, y_train)
6
7 # Make predictions using the testing set
8 y_pred = model.predict(X_test)
9
10 # Creating the DataFrame with features and their coefficients
11 coeff_df = pd.DataFrame({'Feature': X_train.columns, 'Coefficient': model.coef_})
12
13 # Install the tabulate module
14 !pip install tabulate
15
16 # Import the tabulate function
17 from tabulate import tabulate
18
19 # Displaying the DataFrame as a table
20 print("Intercept:", model.intercept_)
21 print("Estimated features vs coefficients:")
22 print(tabulate(coeff_df, headers='keys', tablefmt='pretty'))
```

```
Requirement already satisfied: tabulate in /usr/local/lib/python3.10/dist-packages (0.9.0)
Intercept: 13.238798996804732
Estimated features vs coefficients:
+-----+
| | Feature | Coefficient |
+-----+
| 0 | T_atm | -0.055535893239226766 |
| 1 | Humidity | 0.003422392005267033 |
| 2 | Distance | 0.0026393537872089093 |
| 3 | T_FH_Max1 | 0.7092038079633628 |
| 4 | Age_18-20 | -0.09521023788460654 |
| 5 | Age_21-25 | -0.10911840778186253 |
| 6 | Age_21-30 | -0.03372839990148858 |
| 7 | Age_26-30 | -0.08394109824533079 |
| 8 | Age_31-40 | -0.13698172540270473 |
| 9 | Age_41-50 | 0.21089662330066358 |
| 10 | Age_51-60 | -0.08176010339345889 |
| 11 | Age_>60 | 0.32984334930878834 |
+-----+
```

Figure 11: Output

```
1 plt.figure(figsize=(10, 6))
2 plt.scatter(y_test, y_pred, alpha=0.6, color='blue')
3 plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'k--', lw=2, color='
  red')
4 plt.xlabel('Actual Values')
5 plt.ylabel('Predicted Values')
6 plt.title('Actual vs. Predicted Values')
7 plt.show()
```

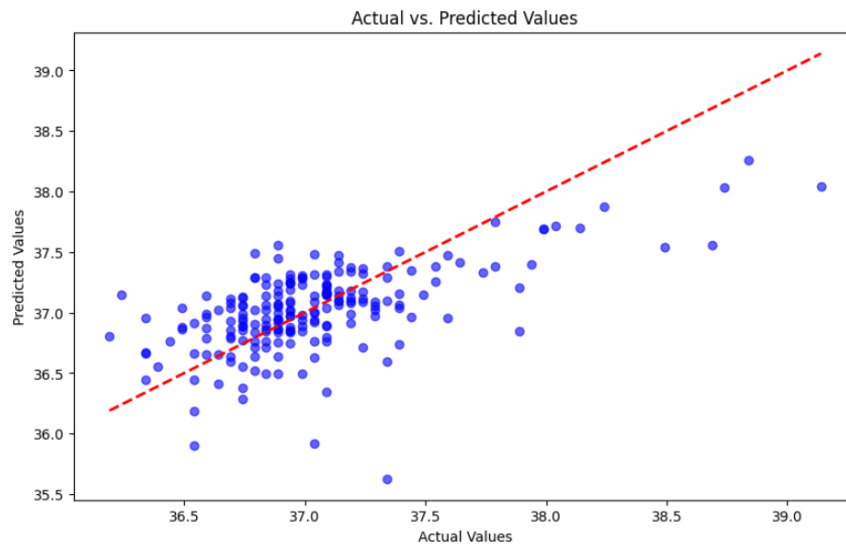


Figure 12: Output

### 3.8 Question 8 : Highly contributes independent variable.

```

1 # Plotting the coefficients
2 !pip install seaborn # install seaborn
3 import seaborn as sns
4 plt.figure(figsize=(10, 6))
5 sns.barplot(x='Coefficient', y='Feature', data=coeff_df, palette='viridis')
6 plt.title('Feature Coefficients in Linear Regression Model')
7 plt.xlabel('Coefficient Value')
8 plt.ylabel('Feature')
9 plt.show()

```

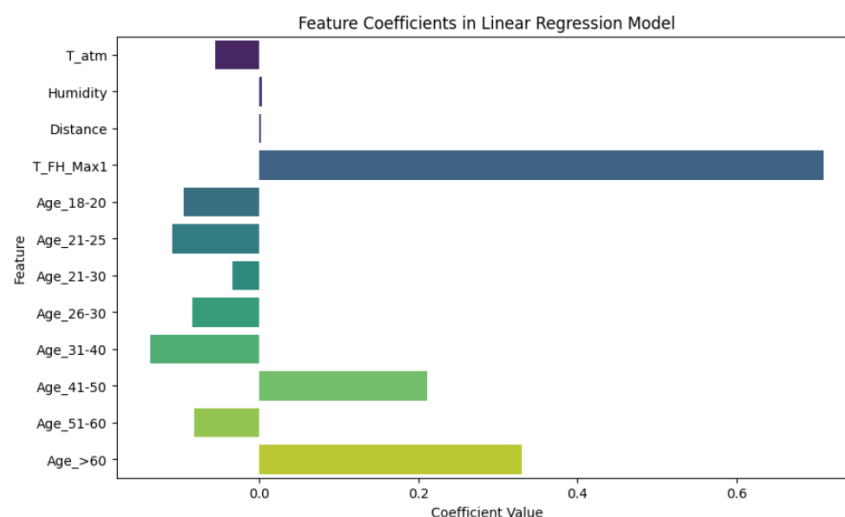


Figure 13: Output

By looking at the coefficient we can see T\_FH\_Max1 got the maximum value which means T\_FH\_Max1 contributes highly for the dependent feature

### 3.9 Question 9 : Select 'T\_OR1', 'T\_OR\_Max1', 'T\_FHC\_Max1', 'T\_FH\_Max1' features as independent features

```
1 # Select the dependent and independent variables
2 dependent_feature_n = data_set['aveOralm']
3 independent_features_n = data_set[['T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']]
4 # Combine into a new DataFrame for modeling
5 model_data_n = pd.concat([independent_features_n, dependent_feature_n], axis=1)
6
7 # Display the prepared dataset
8 print("Prepared dataset:")
9 independent_features_n.head()
10
11 X_train, X_test, y_train, y_test = train_test_split(independent_features_n,
12                                                     dependent_feature_n, test_size=0.2, random_state=42)
13
14 from sklearn import datasets, linear_model
15 from sklearn.metrics import mean_squared_error, r2_score
16
17 model = LinearRegression()
18 model.fit(X_train, y_train)
19
20 # Make predictions using the testing set
21 y_pred = model.predict(X_test)
22
23 # Creating the DataFrame with features and their coefficients
24 coeff_df = pd.DataFrame({'Feature': X_train.columns, 'Coefficient': model.coef_})
25
26 # Displaying the DataFrame as a table
27 print("Intercept:", model.intercept_)
28 print("Estimated features vs coefficients:")
29 print(tabulate(coeff_df, headers='keys', tablefmt='pretty'))
```

```
Prepared dataset:
Intercept: 6.79355629984887
Estimated features vs coefficients:
+---+-----+-----+
|   | Feature | Coefficient |
+---+-----+-----+
| 0 | T_OR1   | 0.20545776323994563 |
| 1 | T_OR_Max1 | 0.34819684316002775 |
| 2 | T_FHC_Max1 | -0.08371846705362093 |
| 3 | T_FH_Max1 | 0.376564342065323 |
+---+-----+-----+
```

Figure 14: Output

```

1 plt.figure(figsize=(10, 6))
2 plt.scatter(y_test, y_pred, alpha=0.6, color='blue')
3 plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'k--', lw=2, color='
  red')
4 plt.xlabel('Actual Values')
5 plt.ylabel('Predicted Values')
6 plt.title('Actual vs. Predicted Values')
7 plt.show()

```

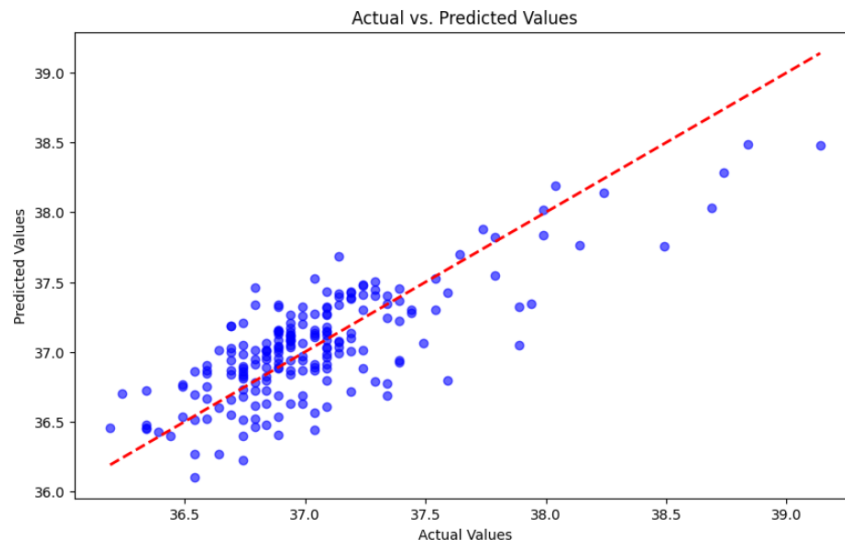


Figure 15: Output

### 3.10 Question 9 : Calculations

```

1 yhat=model.predict(X_train)
2 d = independent_features_n.shape[1]
3 print('d:',d)
4
5 # Residual Sum of Sqaures (RSS)
6 RSS = np.sum((yhat - y_train)**2)
7 print('RSS=', RSS)
8
9 N=len(y_train)
10 #print('Number of Datapoints=',N)
11
12 # Residual Standard Error (RSE)
13 RSE = np.sqrt(1/(N-d-1)*RSS)
14 print('RSE=', RSE)
15
16 # Mean squre error(MSE)
17 #predictions = lm.predict(X_train)
18
19 newX = pd.DataFrame({"Constant":np.ones(len(X_train))}).join(pd.DataFrame(X_train))
20 MSE = (sum((y_train-yhat)**2))/(len(newX)-len(newX.columns))
21 print('MSE=', MSE)
22
23 # Total Sum of Squares (TSS)
24 TSS = np.sum((y_train- np.mean(y_train))**2)
25 print('TSS=', TSS)
26

```

```

27 # R2
28
29 R2 = (TSS - RSS)/TSS
30 print('R (from direct calculations)=', R2)
31
32 # Calculation of R2 using sklearn
33 R2 = model.score(X_train,y_train)
34
35 print('R (from sklearn module)=', R2)

```

RSS= 77.97449082857881  
 RSE= 0.31045739777400483  
 MSE= 0.09638379583260669  
 TSS= 223.63911855036855  
 R<sup>2</sup> (from direct calculations)= 0.6513378726673116  
 R<sup>2</sup> (from sklearn module)= 0.6513378726673116

Calculate the standard error for each feature , t-statistic for each feature and p-value for each feature

```

1 from scipy.stats import t
2 import numpy as np
3 import scipy.stats as stats
4 #samples size
5 SN = len(X_train)
6
7 features = ['T_OR1', 'T_OR_Max1', 'T_FHC_Max1', 'T_FH_Max1']
8 w_0 = model.intercept_
9 w_1 = model.coef_
10
11 NF = len(features)
12
13 #Calculate Standard error
14 standard_error = []
15 for feature in features:
16     SE2 = RSS/(SN-NF-1) / np.sum((X_train[feature] - np.mean(X_train[feature]))**2)
17     standard_error.append(np.sqrt(SE2))
18
19 #calculate t values
20 t_values = []
21 for i in range(NF):
22     t_values.append(w_1[i]/standard_error[i])
23
24 #calculate p-values
25 p_values = []
26 for i in range(NF):
27     p_values.append(2 * (1 - stats.t.cdf(abs(t_values[i]), SN-NF-1)))
28
29
30 # Create a DataFrame to display the values in a table
31 results_df = pd.DataFrame({
32     'Feature': features,
33     'Standard Error': standard_error,
34     't-value': t_values,
35     'p-value': p_values
36 })
37
38 # Print the table
39 print(results_df)
40
41 # Save the table to a CSV file (optional)
42 results_df.to_csv('model_statistics.csv', index=False)

```

	Feature	Standard Error	t-value	p-value
0	T_OR1	0.019226	10.686208	0.00000
1	T_OR_Max1	0.019219	18.117535	0.00000
2	T_FHC_Max1	0.018810	-4.450810	0.00001
3	T_FH_Max1	0.020771	18.129689	0.00000

Figure 16: Output

### 3.11 Question 11

The p-values are very low, which shows there's strong evidence against the null hypothesis. This means we can confidently reject the null hypothesis.

## 4. Performance evaluation of Linear regression

### 4.1 Question 2 : Residual standard error (RSE) for models A and B

The Residual Standard Error (RSE) formula is given by:

$$RSE = \sqrt{\frac{RSS}{N - d - 1}}$$

For Model A:

$$RSE_A = \sqrt{\frac{9}{10^4 - 2 - 1}} = 0.03$$

For Model B:

$$RSE_B = \sqrt{\frac{2}{10^4 - 4 - 1}} = 0.0141$$

Since  $RSE_B$  is less than  $RSE_A$ , Model B shows a better fit compared to Model A.

### 4.2 Question 3 : R-squared (R<sup>2</sup>) for models A and B

The  $R^2$  metric is computed as:

$$R^2 = 1 - \frac{SSE}{TSS}$$

For Model A:

$$R_A^2 = 1 - \frac{9}{90} = 0.9$$

For Model B:

$$R_B^2 = 1 - \frac{2}{10} = 0.8$$

Based on the  $R^2$  values, since  $R_A^2$  is closer to 1 than  $R_B^2$ , Model A performs better in terms of variance explained by the model.

### 4.3 Question 4 : Between RSE and R-squared which performance metric is more fair

The  $R^2$  value is generally preferred when comparing models because it is independent of the scale of  $y$ , making it a more consistent measure across different datasets.



## 5. Impact of Outliers on Linear Regression

### 5.1 Question2 : when $a \rightarrow 0$

Let's define the loss functions  $L_1(\omega)$  and  $L_2(\omega)$ :

$$L_1(\omega) = \frac{1}{N} \sum_{i=1}^N \left( \frac{r_i^2}{a^2 + r_i^2} \right)$$

As  $a \rightarrow 0$ :

$$\lim_{a \rightarrow 0} L_1(\omega) = \frac{1}{N} \sum_{i=1}^N 1 = 1$$

Now, for  $L_2(\omega)$ :

$$L_2(\omega) = \frac{1}{N} \sum_{i=1}^N \left( 1 - e^{-\frac{2|r_i|}{a}} \right)$$

As  $a \rightarrow 0$ :

$$\lim_{a \rightarrow 0} L_2(\omega) = \frac{1}{N} \sum_{i=1}^N 1 = 1$$

Both  $L_1(\omega)$  and  $L_2(\omega)$  tend to 1 as  $a$  approaches 0, meaning they both become constant under this condition.

### 5.2 Question 3 : To minimize the influence of data points

When reducing the effect of outliers, such as for residuals  $|r_i| \geq 40$ , the loss function should be chosen based on how quickly it reduces the influence of these large residuals.

**Comparing  $L_1(\omega)$  and  $L_2(\omega)$ :**

**For  $L_1(\omega)$ :**

$$L_1(\omega) = \frac{1}{N} \sum_{i=1}^N \frac{r_i^2}{a^2 + r_i^2}$$

This function diminishes the impact of outliers gradually. When the residuals are small,  $L_1(\omega)$  behaves similarly to a squared loss. However, as the residuals grow, the function saturates, reducing the influence of larger residuals but not as quickly as  $L_2(\omega)$ .

**For  $L_2(\omega)$ :**

$$L_2(\omega) = \frac{1}{N} \sum_{i=1}^N \left( 1 - e^{-\frac{2|r_i|}{a}} \right)$$

This loss function reduces the impact of large residuals more rapidly due to its exponential term. For smaller residuals, it behaves similarly to a linear loss function, but for larger residuals, it reaches saturation much faster.

**Conclusion:** Based on the plot, it is evident that  $L_2(\omega)$  approaches saturation faster than  $L_1(\omega)$ , making it more suitable for handling outliers. By choosing  $a = 40$ , the loss function ensures that residuals with  $|r_i| \geq 40$  have minimal impact on the model.